

# MDM-2025 Homework 4

Mert Dagli (100830277)  
Tuomas Mäkinen (103382573)  
Maja Sellmer (103396682)

November 24, 2025

# 1 Methods

The task was to explore associations between molecule substructures and anticancer activities using frequent subgraph mining. To this end, we randomly selected one molecule(number 14) as a test sample and then mined frequent subgraphs of the rest molecules using the software *MoSS*. We searched for significant rules and used them to predict anticancer activities of the test molecule. We completed the assignment in Python, using the following libraries:

- `matplotlib` for plotting
- `pandas` to handle the given data as a dataframe
- `math` for calculations
- `sklearn.metrics` for mutual information score
- `rdkit.Chem` for drawing molecules

## 1.1 Frequent subgraph mining

Firstly, we identified substructures in the molecules with the *MoSS* software. In order to obtain interesting frequent substructures, we set the minimum size to 6 and minimum support to 12. Then we identified how common each substructure is. We computed the frequency of a structure as the number of molecules containing this substructure divided by the total number of molecules, and visualised the top 5 most frequent ones.

## 1.2 Significant rule search

Next, for each of the three types of cancer, we calculated the mutual information for rules of the form  $X \rightarrow C$ , where  $X$  represents the co-occurrence of some molecular substructures and  $C$  refers to the positive anticancer activity against that cancer. We only considered sets  $X$  with at most two elements, i.e. single substructures or pairs of two substructures. We only keep rules which are positively associated with anticancer activity, i.e.  $P(C = c | X) > P(C = c)$  and kept those with mutual information higher than the threshold of 0.01.

## 1.3 Selection of representative rules

After identifying the significant rules, we tried to remove overfitting rules. We selected a representative rule set using confidence (rule  $X \rightarrow C$  is removed if there exists  $Y \subset X$  such that  $P(C = c | Y) \geq P(C = c | X)$ ) and conditional mutual

information (a rule is removed if the conditional mutual information falls under the threshold of 0.01).

## 1.4 Prediction

Lastly, we checked which substructures our test molecule contains and thus which rules apply. We computed the average confidence, average MI, average MI\*confidence, average lift, and average leverage of these rules. Based on these goodness measures, we predict whether there is an anticancer activity against each cancer type.

## 2 Results

### 2.1 Frequent substructures

Figure 1 shows the give most frequent molecular substructures.

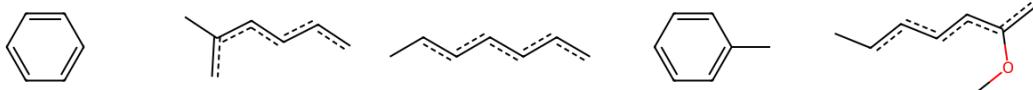


Figure 1: Top 5 substructures

### 2.2 Significant and representative rules

Table 1 shows a summary of the rules we obtained for each cancer type. We can see that the pool of possible rules was significantly narrowed down for types 2 and 3, but it seems the threshold was too low for type 1.

Type, Rules	# rules	Avg freq	Avg conf	Avg MI
Type 1, Significant	1680	0.242	0.896	0.328
Type 1, Representative	482	0.275	0.957	0.489
Type 2, Significant	1475	0.122	0.947	0.138
Type 2, Representative	91	0.198	0.910	0.185
Type 3, Significant	1687	0.115	0.992	0.073
Type 3, Representative	60	0.214	0.976	0.130

Table 1: Summary statistics of significant and representative rules

## 2.3 Predictions

Table 2 shows the average of each goodness measure for the rules applicable to the test molecule. We can see that type 1 seems very likely, although it is not true. If we base our predictions on the lift and leverage, we would incorrectly predict 1 to be true, correctly predict 2 to be true, and incorrectly predict 3 to be false. We could classify correctly by using a confidence threshold of 0.915, but this would have had to be a very lucky guess. It seems the average confidence of the rules we found is overall too high.

Type	Confidence	MI	MI × Conf	Lift	Leverage	True?
Type 1	0.913	0.337	0.314	2.624	0.565	0
Type 2	0.918	0.051	0.045	2.483	0.548	1
Type 3	0.993	0.079	0.078	1.522	0.341	1

Table 2: Predictions of anticancer activities for the test molecule

## 3 Conclusions

To summarize, we identified frequent substructures in the molecules, and searched for rules telling us that the presence of some substructure  $X$ , or a pair of substructures  $(X_1, X_2)$  implies anticancer activity against cancer type  $n$ . We narrowed down possible rules using confidence, mutual information and conditional mutual information. This allowed us to make predictions about anticancer activites of a test molecule not included in the training data, however we were not really able to predict all three types correctly based on the rules we obtained.

## 4 Appendix – Code

```
import os, sys, random, itertools, math
import pandas as pd
import numpy as np
from sklearn.metrics import mutual_info_score
from rdkit import Chem
from rdkit.Chem import Draw
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import ast

#-----Prep-----#
MOLECULE_CSV = 'molecule.csv'
MOSS_IDENTIFIERS = f"newdata/moss_identifiers.csv"
MOSS_SUBSTRUCTURES = f"newdata/moss_substructures.csv"
molecule_df = pd.read_csv(MOLECULE_CSV, sep=";")
```

```

# Excluded Test Molecule
test_mol = molecule_df.iloc[14]
molecule_df = molecule_df.drop([14])

with open("molecule_moss.smile", "w") as f:
    for i, row in molecule_df.iterrows():
        mol_id = f"{i+1}"
        smiles = row["SMILES"]
        f.write(f"{mol_id},0,{smiles}\n")

#-----A-----
molecule_df = pd.read_csv(MOLECULE_CSV, sep=";")
# Excluded Molecule
test_mol = molecule_df.iloc[14]
molecule_df = molecule_df.drop([14])

moss_df = pd.read_csv(MOSS_SUBSTRUCTURES, sep=",")
moss_df["mol_ids"] = pd.read_csv(MOSS_IDENTIFIER, sep=";")['list'].apply(lambda x
    : [int(i) for i in ast.literal_eval(x)])

cleaned_moss_df = moss_df.drop_duplicates(subset=['description']).copy()
    .reset_index()
cleaned_moss_df.rename(columns={'description': 'substructure'}, inplace=True)
cleaned_moss_df["frq"] = (cleaned_moss_df["mol_ids"].apply(len) / molecule_df.
    shape[0])

print(f"Original size: {moss_df.shape[0]}")
print(f"Cleaned size: {cleaned_moss_df.shape[0]}")
print(f"Removed {len(moss_df) - len(cleaned_moss_df)} duplicates") #NOTE! Idk
    why there are duplicates. Use the cleaned set. (type 3 is the worst, others
    don't lose much)

sorted_df = cleaned_moss_df.sort_values("frq", ascending=False).head(5).drop(
    columns=["nodes", "edges", "s_abs", "s_rel", "c_abs", "c_rel", "id"])
top5 = sorted_df.head(5)
top5_mols = []
test = [0,0,0,0,0]
for sml in top5["substructure"]:
    mol = Chem.MolFromSmiles(sml, sanitize=False)
    top5_mols.append(mol)

for sml in molecule_df["SMILES"]:
    mol = Chem.MolFromSmiles(sml)
    for i, q_mol in enumerate(top5_mols):
        if mol.HasSubstructMatch(q_mol):
            test[i] += 1

img = Draw.MolsToGridImage(top5_mols, molsPerRow=5, subImgSize=(250,250),
    returnPNG=False)
img.save(f"results/top5_substructures.png")
print(test)

#-----B-----
molecule_df["id"] = list(molecule_df.index.values + 1)
mis = []

def mutual_information_binary(X, C):
    X = np.asarray(X)

```

```

C = np.asarray(C)

P11 = np.mean((X == 1) & (C == 1))
P10 = np.mean((X == 1) & (C == 0))
P01 = np.mean((X == 0) & (C == 1))
P00 = np.mean((X == 0) & (C == 0))

PX1 = P11 + P10
PX0 = P01 + P00
PC1 = P11 + P01
PC0 = P10 + P00

def calc(Pxy, Px, Pc):
    if Pxy == 0 or Px == 0 or Pc == 0:
        return 0
    return Pxy * np.log2(Pxy / (Px * Pc))

MI = (calc(P11, PX1, PC1) + calc(P10, PX1, PC0) + calc(P01, PX0, PC1) + calc(P00, PX0, PC0))

return MI

# CHOOSE CANCER TYPE n = 1, 2, 3
n = 3

MI_threshold = 0.01
MIC_threshold = 0.01

molecule_df["id"] = list(molecule_df.index.values + 1)

significant_MI_ids = []
significant_mis = []
significant_supports = []
significant_P_C_given_X = []

rep_MI_ids = []
rep_mis = []
rep_supports = []
rep_P_C_given_X = []

C = molecule_df[f'anti_cancer_{n}'].values
P_C_global = C.mean()

for idx1, row1 in cleaned_moss_df.iterrows():
    for idx2, row2 in cleaned_moss_df.iterrows():
        if idx1 > idx2:
            continue
        if idx1 == idx2:
            pair_id = f"{row1['id']}"
        else:
            pair_id = f"{row1['id']}&{row2['id']}"

        X1 = molecule_df["id"].isin(row1['mol_ids']).values
        X2 = molecule_df["id"].isin(row2['mol_ids']).values
        X = X1 & X2

        if np.sum(X) == 0:
            continue

        P_C_X = C[X].mean()
        if P_C_X <= P_C_global:
            continue

        if pair_id in significant_MI_ids:
            rep_MI_ids.append(pair_id)
        else:
            significant_MI_ids.append(pair_id)

        if pair_id in significant_mis:
            rep_mis.append(pair_id)
        else:
            significant_mis.append(pair_id)

        if pair_id in significant_supports:
            rep_supports.append(pair_id)
        else:
            significant_supports.append(pair_id)

        if pair_id in significant_P_C_given_X:
            rep_P_C_given_X.append(pair_id)
        else:
            significant_P_C_given_X.append(pair_id)

```

```

MI = mutual_information_binary(X, C)
if MI < MI_threshold:
    continue

significant_MI_ids.append(pair_id)
significant_mis.append(MI)
significant_supports.append(X.mean())
significant_P_C_given_X.append(P_C_X)

if idx1 != idx2:
    P_C_X1 = C[X1].mean()
    P_C_X2 = C[X2].mean()

    if P_C_X <= max(P_C_X1, P_C_X2):
        continue

    MI_Y1 = mutual_information_binary(X1, C)
    MI_Y2 = mutual_information_binary(X2, C)
    MIC = MI - max(MI_Y1, MI_Y2)
    if MIC <= MIC_threshold:
        continue

    rep_MI_ids.append(pair_id)
    rep_mis.append(MI)
    rep_supports.append(X.mean())
    rep_P_C_given_X.append(P_C_X)

sig_df = pd.DataFrame({
    "id": significant_MI_ids,
    f"MI_anti_cancer_{n)": significant_mis,
    "support": significant_supports,
    "P_C_given_X": significant_P_C_given_X
})

rep_df = pd.DataFrame({
    "id": rep_MI_ids,
    f"MI_anti_cancer_{n)": rep_mis,
    "support": rep_supports,
    "P_C_given_X": rep_P_C_given_X
})

print("Significant rules summary")
print(sig_df.head())
stats = {
    'num_rules': len(sig_df),
    'avg_support': sig_df['support'].mean(),
    'avg_confidence': sig_df['P_C_given_X'].mean(),
    'avg_MI': sig_df[f'MI_anti_cancer_{n}'].mean(),
}
print(stats)

print("Representative rules summary")
print(rep_df.head())
stats = {
    'num_rules': len(rep_df),
    'avg_support': rep_df['support'].mean(),
    'avg_confidence': rep_df['P_C_given_X'].mean(),
    'avg_MI': rep_df[f'MI_anti_cancer_{n}'].mean(),
}
print(stats)

```

```

mol = Chem.MolFromSmiles(test_mol["SMILES"])
test_substructures = []
for i,q_mol_smile in enumerate(cleaned_moss_df["substructure"]):
    q_mol = Chem.MolFromSmiles(q_mol_smile, sanitize=False)
    if mol.HasSubstructMatch(q_mol):
        test_substructures.append(i +1)
print(test_substructures)

selected_rules = []

for idx, row in rep_df.iterrows():
    rule_id = row['id']
    sub_ids = list(map(int, rule_id.split('&')))
    if all(s in test_substructures for s in sub_ids):
        selected_rules.append(row)

selected_rules_df = pd.DataFrame(selected_rules)
print(selected_rules_df)

C = molecule_df[f'anti_cancer_{n}'].values
P_C_global = np.mean(C)

selected_rules_df['MI_conf'] = selected_rules_df[f'MI_anti_cancer_{n}'] *
    selected_rules_df['P_C_given_X']
selected_rules_df['lift'] = selected_rules_df['P_C_given_X'] / P_C_global
selected_rules_df['leverage'] = selected_rules_df['P_C_given_X'] - P_C_global

stats = {
    'avg_confidence': selected_rules_df['P_C_given_X'].mean(),
    'avg_MI': selected_rules_df[f'MI_anti_cancer_{n}'].mean(),
    'avg_MI_conf': selected_rules_df['MI_conf'].mean(),
    'avg_lift': selected_rules_df['lift'].mean(),
    'avg_leverage': selected_rules_df['leverage'].mean()
}
print(stats)

predicted_activity = 1 if stats['avg_confidence'] > 0.85 else 0
true_activity = test_mol[f'anti_cancer_{n}']
print(f"Predicted: {predicted_activity}, True: {true_activity}")

```

## 5 Bibliography

### References

- [1] Charu C. Aggarwal. *Data Mining - The Textbook*. Springer, 2015.